

## PROJECT-3 REPORT (NAVJOT KAUR: 301404765)

**Kaggle username: Navjott77**

I am claiming my 3 late days for this assignment

## PART 1: OBJECT DETECTION

In this part of the project the train dataset is loaded from train.json and the steps for loading test dataset by keeping the annotations empty were followed as listed in the project description.

Images from train dataset to make sure the function works correctly:





### 1. List of configs and modifications that I used:

```
cfg = get_cfg()
cfg.OUTPUT_DIR = "{}output/".format(BASE_DIR)
cfg.merge_from_file(model_zoo.get_config_file("COCO-
Detection/faster_rcnn_R_101_FPN_3x.yaml"))
cfg.DATASETS.TRAIN = ("plane_train",)
cfg.DATASETS.TEST = ()
cfg.DATALOADER.NUM_WORKERS = 2
cfg.SOLVER.IMS_PER_BATCH = 2
cfg.SOLVER.BASE_LR = 0.00025
cfg.SOLVER.MAX_ITER = 850
cfg.SOLVER.STEPS = []
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 512
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 1
```

### 2. Factors which helped improve the performance. Explain each factor in 2-3 lines.

The provided configurations in the description were tested first where only the max iterations is changed from 500 to 850 to minimize the total\_loss. At the end the loss became as low as 1.02 starting from approximate 1.98, which was quite good.

Base Learning Rate was kept as 0.00025 after using smaller and larger rates but the results responded better when the rate was 0.00025. When the LR was larger the accuracy was much low and when it was smaller then it started off with a good accuracy but remained constant over time.

Other factors such as num\_workers, ims\_per\_batch, batch\_size\_per\_image and num\_classes were kept constant because they did not have much effect on the model.



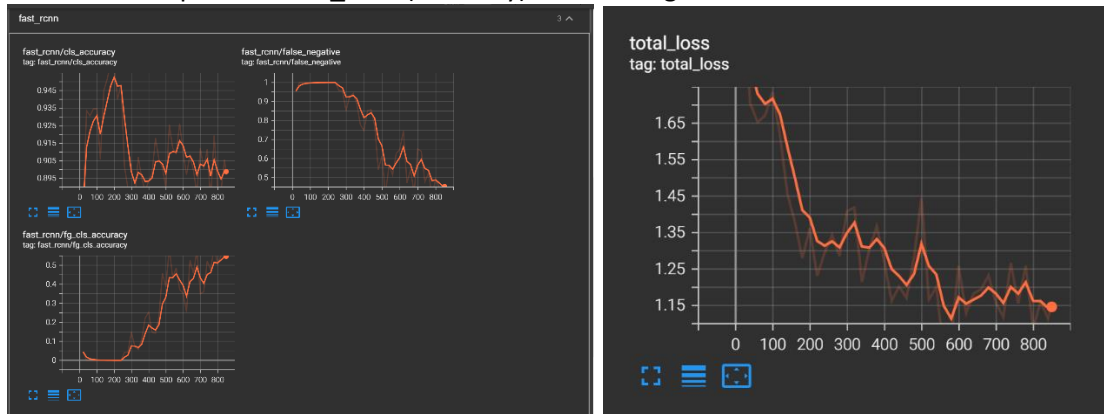
```

8 d2.engine.train_loop: Starting training from iteration 0
3 d2.utils.events]: eta: 0:13:51 iter: 19 total_loss: 1.989 loss_cls: 0.4796 loss_box_reg: 0.04301 loss_rpn_cls: 0.6727 loss_rpn_loc: 0.7133 time: 1.143
8 d2.utils.events]: eta: 0:13:32 iter: 39 total_loss: 1.704 loss_cls: 0.3004 loss_box_reg: 0.02583 loss_rpn_cls: 0.6791 loss_rpn_loc: 0.7375 time: 1.209
0 d2.utils.events]: eta: 0:11:37 iter: 59 total_loss: 1.653 loss_cls: 0.2651 loss_box_reg: 0.05806 loss_rpn_cls: 0.669 loss_rpn_loc: 0.6594 time: 1.1657
6 d2.utils.events]: eta: 0:11:47 iter: 79 total_loss: 1.671 loss_cls: 0.2472 loss_box_reg: 0.02023 loss_rpn_cls: 0.6685 loss_rpn_loc: 0.6038 time: 1.206
1 d2.utils.events]: eta: 0:11:50 iter: 99 total_loss: 1.735 loss_cls: 0.233 loss_box_reg: 0.02978 loss_rpn_cls: 0.6561 loss_rpn_loc: 0.6485 time: 1.2071
7 d2.utils.events]: eta: 0:11:48 iter: 119 total_loss: 1.618 loss_cls: 0.2931 loss_box_reg: 0.04188 loss_rpn_cls: 0.6158 loss_rpn_loc: 0.6087 time: 1.22
0 d2.utils.events]: eta: 0:11:50 iter: 139 total_loss: 1.456 loss_cls: 0.2178 loss_box_reg: 0.02754 loss_rpn_cls: 0.6193 loss_rpn_loc: 0.5675 time: 1.21
2 d2.utils.events]: eta: 0:11:10 iter: 159 total_loss: 1.379 loss_cls: 0.1718 loss_box_reg: 0.03791 loss_rpn_cls: 0.5909 loss_rpn_loc: 0.5754 time: 1.19
4 d2.utils.events]: eta: 0:11:02 iter: 179 total_loss: 1.28 loss_cls: 0.1642 loss_box_reg: 0.03669 loss_rpn_cls: 0.5875 loss_rpn_loc: 0.4436 time: 1.188
0 d2.utils.events]: eta: 0:10:31 iter: 199 total_loss: 1.361 loss_cls: 0.1576 loss_box_reg: 0.01372 loss_rpn_cls: 0.5567 loss_rpn_loc: 0.5021 time: 1.19
6 d2.utils.events]: eta: 0:10:33 iter: 219 total_loss: 1.231 loss_cls: 0.1802 loss_box_reg: 0.02521 loss_rpn_cls: 0.526 loss_rpn_loc: 0.4713 time: 1.209
9 d2.utils.events]: eta: 0:10:02 iter: 239 total_loss: 1.295 loss_cls: 0.1912 loss_box_reg: 0.0622 loss_rpn_cls: 0.5258 loss_rpn_loc: 0.3793 time: 1.204
1 d2.utils.events]: eta: 0:09:26 iter: 259 total_loss: 1.344 loss_cls: 0.2541 loss_box_reg: 0.1457 loss_rpn_cls: 0.5246 loss_rpn_loc: 0.3383 time: 1.199
3 d2.utils.events]: eta: 0:09:23 iter: 279 total_loss: 1.285 loss_cls: 0.26 loss_box_reg: 0.1851 loss_rpn_cls: 0.4763 loss_rpn_loc: 0.3601 time: 1.1880
0 d2.utils.events]: eta: 0:09:16 iter: 299 total_loss: 1.409 loss_cls: 0.2849 loss_box_reg: 0.2663 loss_rpn_cls: 0.47 loss_rpn_loc: 0.4357 time: 1.1999
2 d2.utils.events]: eta: 0:08:53 iter: 319 total_loss: 1.419 loss_cls: 0.2706 loss_box_reg: 0.2499 loss_rpn_cls: 0.4533 loss_rpn_loc: 0.4148 time: 1.1937
7 d2.utils.events]: eta: 0:08:36 iter: 339 total_loss: 1.214 loss_cls: 0.207 loss_box_reg: 0.1247 loss_rpn_cls: 0.4362 loss_rpn_loc: 0.3025 time: 1.1962
8 d2.utils.events]: eta: 0:08:11 iter: 359 total_loss: 1.304 loss_cls: 0.2583 loss_box_reg: 0.2649 loss_rpn_cls: 0.4322 loss_rpn_loc: 0.3127 time: 1.1887
3 d2.utils.events]: eta: 0:07:51 iter: 379 total_loss: 1.367 loss_cls: 0.2458 loss_box_reg: 0.2758 loss_rpn_cls: 0.3953 loss_rpn_loc: 0.3102 time: 1.1919
8 d2.utils.events]: eta: 0:07:32 iter: 399 total_loss: 1.272 loss_cls: 0.2721 loss_box_reg: 0.2384 loss_rpn_cls: 0.4059 loss_rpn_loc: 0.3034 time: 1.1944
3 d2.utils.events]: eta: 0:07:13 iter: 419 total_loss: 1.163 loss_cls: 0.244 loss_box_reg: 0.2479 loss_rpn_cls: 0.3817 loss_rpn_loc: 0.3307 time: 1.1966
7 d2.utils.events]: eta: 0:06:56 iter: 439 total_loss: 1.282 loss_cls: 0.2128 loss_box_reg: 0.2024 loss_rpn_cls: 0.3551 loss_rpn_loc: 0.3553 time: 1.1959
1 d2.utils.events]: eta: 0:06:34 iter: 459 total_loss: 1.17 loss_cls: 0.2401 loss_box_reg: 0.2799 loss_rpn_cls: 0.3294 loss_rpn_loc: 0.3639 time: 1.1958
4 d2.utils.events]: eta: 0:06:12 iter: 479 total_loss: 1.28 loss_cls: 0.2501 loss_box_reg: 0.2888 loss_rpn_cls: 0.3228 loss_rpn_loc: 0.2843 time: 1.1927
1 d2.utils.events]: eta: 0:05:52 iter: 499 total_loss: 1.444 loss_cls: 0.2733 loss_box_reg: 0.3554 loss_rpn_cls: 0.3615 loss_rpn_loc: 0.3654 time: 1.2003
3 d2.utils.events]: eta: 0:05:31 iter: 519 total_loss: 1.166 loss_cls: 0.1991 loss_box_reg: 0.3024 loss_rpn_cls: 0.3478 loss_rpn_loc: 0.343 time: 1.1980
7 d2.utils.events]: eta: 0:05:12 iter: 539 total_loss: 1.201 loss_cls: 0.2205 loss_box_reg: 0.4125 loss_rpn_cls: 0.2776 loss_rpn_loc: 0.3082 time: 1.1965
3 d2.utils.events]: eta: 0:04:51 iter: 559 total_loss: 1.02 loss_cls: 0.2177 loss_box_reg: 0.3416 loss_rpn_cls: 0.2504 loss_rpn_loc: 0.2154 time: 1.2015
4 d2.utils.events]: eta: 0:04:31 iter: 579 total_loss: 1.062 loss_cls: 0.1984 loss_box_reg: 0.3109 loss_rpn_cls: 0.266 loss_rpn_loc: 0.326 time: 1.1959
0 d2.utils.events]: eta: 0:04:11 iter: 599 total_loss: 1.259 loss_cls: 0.2247 loss_box_reg: 0.319 loss_rpn_cls: 0.2865 loss_rpn_loc: 0.3539 time: 1.1989
3 d2.utils.events]: eta: 0:03:52 iter: 619 total_loss: 1.129 loss_cls: 0.2592 loss_box_reg: 0.2354 loss_rpn_cls: 0.2446 loss_rpn_loc: 0.2778 time: 1.1974
0 d2.utils.events]: eta: 0:03:32 iter: 639 total_loss: 1.183 loss_cls: 0.2188 loss_box_reg: 0.3855 loss_rpn_cls: 0.2364 loss_rpn_loc: 0.286 time: 1.2014
2 d2.utils.events]: eta: 0:03:13 iter: 659 total_loss: 1.194 loss_cls: 0.2214 loss_box_reg: 0.3656 loss_rpn_cls: 0.231 loss_rpn_loc: 0.2967 time: 1.1988
7 d2.utils.events]: eta: 0:02:52 iter: 679 total_loss: 1.232 loss_cls: 0.2482 loss_box_reg: 0.4428 loss_rpn_cls: 0.1991 loss_rpn_loc: 0.2291 time: 1.1938
2 d2.utils.events]: eta: 0:02:32 iter: 699 total_loss: 1.158 loss_cls: 0.21 loss_box_reg: 0.3829 loss_rpn_cls: 0.2342 loss_rpn_loc: 0.2702 time: 1.1948
1 d2.utils.events]: eta: 0:02:12 iter: 719 total_loss: 1.118 loss_cls: 0.218 loss_box_reg: 0.3564 loss_rpn_cls: 0.2321 loss_rpn_loc: 0.2577 time: 1.1948
5 d2.utils.events]: eta: 0:01:51 iter: 739 total_loss: 1.266 loss_cls: 0.2343 loss_box_reg: 0.4074 loss_rpn_cls: 0.1711 loss_rpn_loc: 0.2232 time: 1.1956
2 d2.utils.events]: eta: 0:01:31 iter: 759 total_loss: 1.156 loss_cls: 0.2081 loss_box_reg: 0.2838 loss_rpn_cls: 0.246 loss_rpn_loc: 0.2995 time: 1.1985
4 d2.utils.events]: eta: 0:01:10 iter: 779 total_loss: 1.26 loss_cls: 0.2055 loss_box_reg: 0.4206 loss_rpn_cls: 0.2209 loss_rpn_loc: 0.2877 time: 1.1968
1 d2.utils.events]: eta: 0:00:50 iter: 799 total_loss: 1.086 loss_cls: 0.2485 loss_box_reg: 0.428 loss_rpn_cls: 0.1921 loss_rpn_loc: 0.2428 time: 1.1997
2 d2.utils.events]: eta: 0:00:30 iter: 819 total_loss: 1.161 loss_cls: 0.2586 loss_box_reg: 0.4602 loss_rpn_cls: 0.1788 loss_rpn_loc: 0.2605 time: 1.1967
5 d2.utils.events]: eta: 0:00:10 iter: 839 total_loss: 1.114 loss_cls: 0.2279 loss_box_reg: 0.3929 loss_rpn_cls: 0.1791 loss_rpn_loc: 0.2728 time: 1.1954
8 d2.utils.events]: eta: 0:00:00 iter: 849 total_loss: 1.151 loss_cls: 0.2334 loss_box_reg: 0.3865 loss_rpn_cls: 0.1961 loss_rpn_loc: 0.2742 time: 1.1947
8 d2.engine.hooks]: Overall training speed: 848 iterations in 0:16:53 (1.1948 s / it)
8 d2.engine.hooks]: Total training time: 0:16:55 (0:00:02 on hooks)


```



### 3. Final plot for total training loss and accuracy

Below are the plots for fast\_rcnn (accuracy) and training loss:



### 4. Visualization of 3 samples from the test set and the predicted results






## 5. Ablation study

The improvements made to learning rate and max\_iter were found to improve the performance of the model. Below are the images that show the difference.

Above 3 images are when iterations were 850.

Below are the images when iterations were 400:







## **PART 2: SEMANTIC SEGMENTATION**

1. Report any hyperparameter settings you used

```
num_epochs = 50
batch_size = 4
learning_rate = 0.001
weight_decay = 1e-5
optimizer = SGD
```

2. Report the final architecture of your network including any modification that you have for the layers. Briefly explain the reason for each modification.

```
#Encoder

self.input_conv = conv(3, 4)
self.down1 = down(4, 16)
```

```

self.down2 = down(16, 32)
self.down3 = down(32, 64)
self.down4 = down(64, 128)
self.down5 = down(128, 256)
self.down6 = down(256, 512)

# Decoder

self.up1 = up(512, 256)
self.up2 = up(256, 128)
self.up3 = up(128, 64)
self.up4 = up(64, 32)
self.up5 = up(32, 16)
self.up6 = up(16, 4)

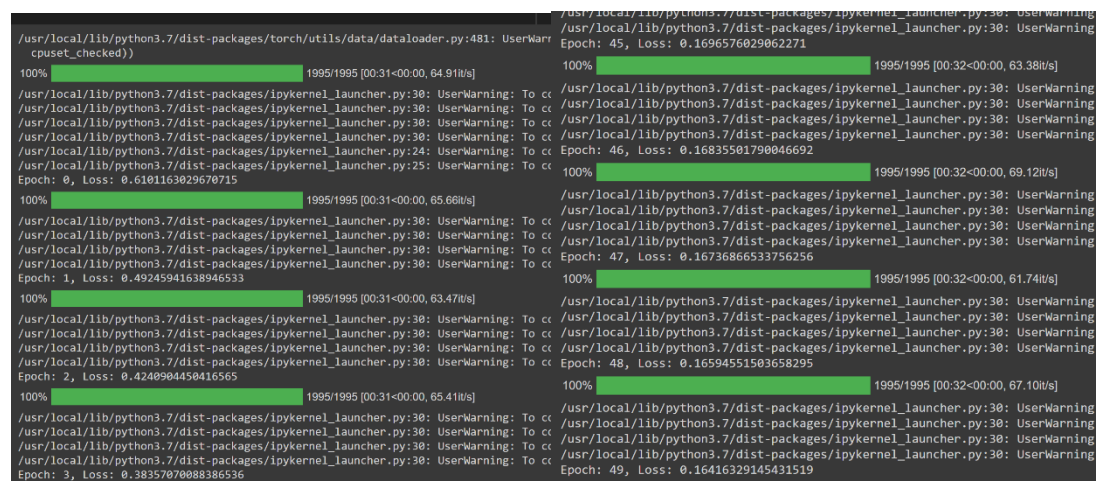
self.output_conv = conv(4, 1, False)

```

The above given architecture is used for the model. The reasoning for the modifications were that firstly increasing the encoder and decoder layers helped in better visualizations and predictions of the plane images. Second the loss reduced drastically from 0.6101 to 0.1641. By keeping the hyperparameters to 50 epochs, batch\_size = 4, learning\_rate as 0.001 and SGD optimizer (given), the feature detector was improved for better detection, visualizations and predictions.

### 3. Report the loss functions that you used and the plot the total training loss of the training procedure

Loss functions used is nn.BCEWithLogitsLoss, which is already given in the code. NO changes were made here. The loss achieved is 0.1641.



```

/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:481: UserWarning:
cpuset_checked))
Epoch: 45, Loss: 0.1696576029062271
100% 1995/1995 [00:31<00:00, 64.91it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:24: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:25: UserWarning:
Epoch: 0, Loss: 0.6101163029670715
100% 1995/1995 [00:31<00:00, 65.66it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
Epoch: 1, Loss: 0.49245941638946533
100% 1995/1995 [00:31<00:00, 63.47it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
Epoch: 2, Loss: 0.4240904450416565
100% 1995/1995 [00:31<00:00, 65.41it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
Epoch: 3, Loss: 0.38357070088386536
Epoch: 46, Loss: 0.16835501790046692
100% 1995/1995 [00:32<00:00, 69.12it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
Epoch: 47, Loss: 0.16736866533756256
100% 1995/1995 [00:32<00:00, 61.74it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
Epoch: 48, Loss: 0.16594551503658295
100% 1995/1995 [00:32<00:00, 67.10it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning:
Epoch: 49, Loss: 0.16416329145431519

```

#### 4. Report the final mean IoU of your model.

IoU is calculated by first calculating the intersection and union as follows and then dividing them.

```
true_img = np rint(np.array(nn.Sigmoid()(pred[i])[0]))

true_mask = np rint(np.array(nn.Sigmoid()(mask[i])[0]))


intersection = np.sum(np.logical_and(true_mask, true_img))
union = np.sum(np.logical_or(true_mask, true_img))
iou = intersection/union

total_iou = total_iou + iou
```

Final IoU of the model is 0.80

```
/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:481: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max
cpuset_checked))
100% 998/998 [00:11<00:00, 87.02it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detac
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detac
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:30: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detac
#images: 7980, Mean IoU: 0.8088094262411628
```

#### 5. Visualize 3 images from the test set and the corresponding predicted masks.



### PART 3: INSTANCE SEGMENTATION

1. The name under which you submitted on Kaggle. Navjott77

2. Report the best score (should match your score on Kaggle).


0.64027

5	Navjott77		0.64027	3	2h
---	-----------	---	---------	---	----

3. The visualisation of results for 3 random samples from the test set.









4. CSV file has been uploaded to the Kaggle.

#### PART 4: Mask R-CNN

1. The visualisation and the evaluation results similar to Part 1.





**2. Explain the differences between the results of Part 1 and Part 4 in a few lines**

The model in Part 1 was better than the one used in Part 4, which is evident from the AP scores and the visualization difference as the planes detected were more in Part 1 than Part 4.

**Part 1: bbox**

AP	AP50	AP75	APs	APm	APl
14.768	37.314	7.014	12.389	21.734	16.720

**Part 4: segm**

AP	AP50	AP75	APs	APm	APl
3.979	15.033	1.173	1.774	4.441	25.165